

# Collaborating & Automation

## Exploring GitLab

# Outline

- Sharing & Collaborating via remotes
- The GitLab flow (Merge Requests)
- Project Management in GitLab
- Forks (Intro)
- Automating with GitLab CI/CD
- GitLab Pages (Intro)
- Auto DevOps

# Remotes

Local references to external repositories

- You'll reference one or more *remotes* from your local repository
- Not a live link / No auto-synchronization

```
# add a remote called "origin"  
git remote add origin git@gitlab.com:rm-training/a.git  
  
# inspect  
git remote -vv  
  
# upload a branch  
git push origin my-branch
```

# GitLab

A platform for the entire software development lifecycle

- Very full-featured
- Open-source, enterprise and on-site available
- Release updates monthly
- Devops focused solutions

*Alternatives include:* GitHub, Bitbucket, Gerrit, etc...

# High level features

- Project Management
- Code Management
  - Code review
  - Common workflow
- Web IDE
- Tight CI/CD Integration
  - Reporting & Monitoring
  - DevOps

## Demo: I'll share my repository

Let's explore GitLab while I push up my personal repository

## To create a new repository...

1. Create an empty repository on the host
2. Add a *remote reference* called `origin`
3. *Pushed* my `master` branch up

```
git remote add <remote name> <url of repo>
git push <remote name> <branch name>
```

I can optionally *track* remote branches

```
git push origin master -u
```

# Project management in GitLab

- Members
- Projects
- Groups

# GitLab Projects

- Projects are repository
- Can have own **Issues, Boards and Milestones**
- Releases map to Git `tag`s
- Merge Requests
- CI/CD settings

# Repository Rule-Setting

- Set access permissions
- Protect special branches
- Require reviews
- Limit who can push to where
- *Auto-sync* across other repositories

## SSH vs HTTPS

You'll likely want to [set up your SSH keys](#) if you haven't already

# Lab: Share your repository

Make sure you have an account!

1. Create a new repository on the host, called "About Me"

- Do not *initialize* with a readme

2. Reference the remote in your local repository

```
git remote add origin <url>
```

3. Push your master branch up, with tracking enabled

```
git push origin master -u
```

4. Visit the project on GitLab! Refresh to see the updates.

5. Locally, list your remotes and all branches

```
git branch -a
```

```
git remote -vv
```

## push to share

Push uploads only the necessary commits reachable in the branch you push

- Typically one branch at a time (but you can push many)
- Keep branch names the same
- It does not matter where you are when you push

```
git push origin bug-1  
git push upstream master
```

```
# send all  
git push origin --all
```

## fetch to refresh local data

Fetch updates your repository with the latest data from the remote

- **It does not merge anything**
  - You'll need to merge manually
- Safe to do from anywhere

```
git fetch origin  
  
# then update your local master  
git checkout master  
git merge origin/master
```

# pull to update a branch

Pull will get the latest data and update the branch you are on

- It does a `fetch` and then `merge` to your current branch
- It does matter where you are - **it affects the branch you are on!**

```
# update master  
git checkout master  
git pull origin master
```

```
# update a branch you are working on  
git checkout 1-my-profile  
git pull origin master
```

```
# update a branch with changes from your team mate  
git checkout 2-our-index  
git pull origin 2-our-index
```

# Remote Branches

Remote branches can be *referenced* locally but they are not really branches for you to work on...

- The *refs* look like `remote-name/branch-name`
- Need to be pruned manually
- Not a branch - more like a "tag"

```
git branch --all  
git branch --remote  
  
git fetch origin --prune  
git push origin --delete 1-feature-work
```

## Want to work on a remote branch?

Just check it out, git will auto-create a local copy and set up tracking, too.

```
$> git branch -a
master
remotes/origin/master
remotes/origin/5-feature-work

$> git checkout 5-feature-work
Branch 'pipeline' set up to track remote branch...
Switched to a new branch 'pipeline'
```

# Tracking Branches

Indicates that a branch is *related* to another branch on the remote

- Does not auto-sync
- Gives you extra info & shortcuts
- Can configure through `branch` or `push`

```
git pull # error! needs tracking  
  
git branch --track 1-bug origin/1-bug  
git pull # success!  
  
# view tracking details of your branches  
git branch -vv
```

## Sharing tags

Tags must be explicitly pushed or configured to do so

```
# push one by one  
git push origin v1  
  
# push all reachable tags  
git push origin --tags  
  
# configure git to always include tags  
git config push.followTags true
```

# Merge Requests

Merge (or Pull) Requests encapsulate a set of changes (in a branch) you want to introduce to the project

- "Please merge my branch"
- Enables discussion, code review & automation around that branch
- Issues can be related to *merge requests*, too
- You can mark it WIP
- aka *Pull Requests* on other platforms

# Quiz

1. How do you share your branch?
2. How do you get your `master` up to date?
3. How do you get your `1-super-feature` branch up to date with the latest from `master` ?
4. What is the difference between `fetch` and `pull`
5. Does it matter what branch you're on?
  - when you `fetch`
  - when you `pull`
  - when you `push`

# Lab: Push & Pull (the GitLab Flow)

On your personal project, make a change and get it merged with a Merge Request, then update your local `master`

## Sharing work

- Make a change in your local repo on a new branch, pushing it up to the remote.

## Integrate your work

- Open a pull/merge request into `master`, review and merge it on GitLab

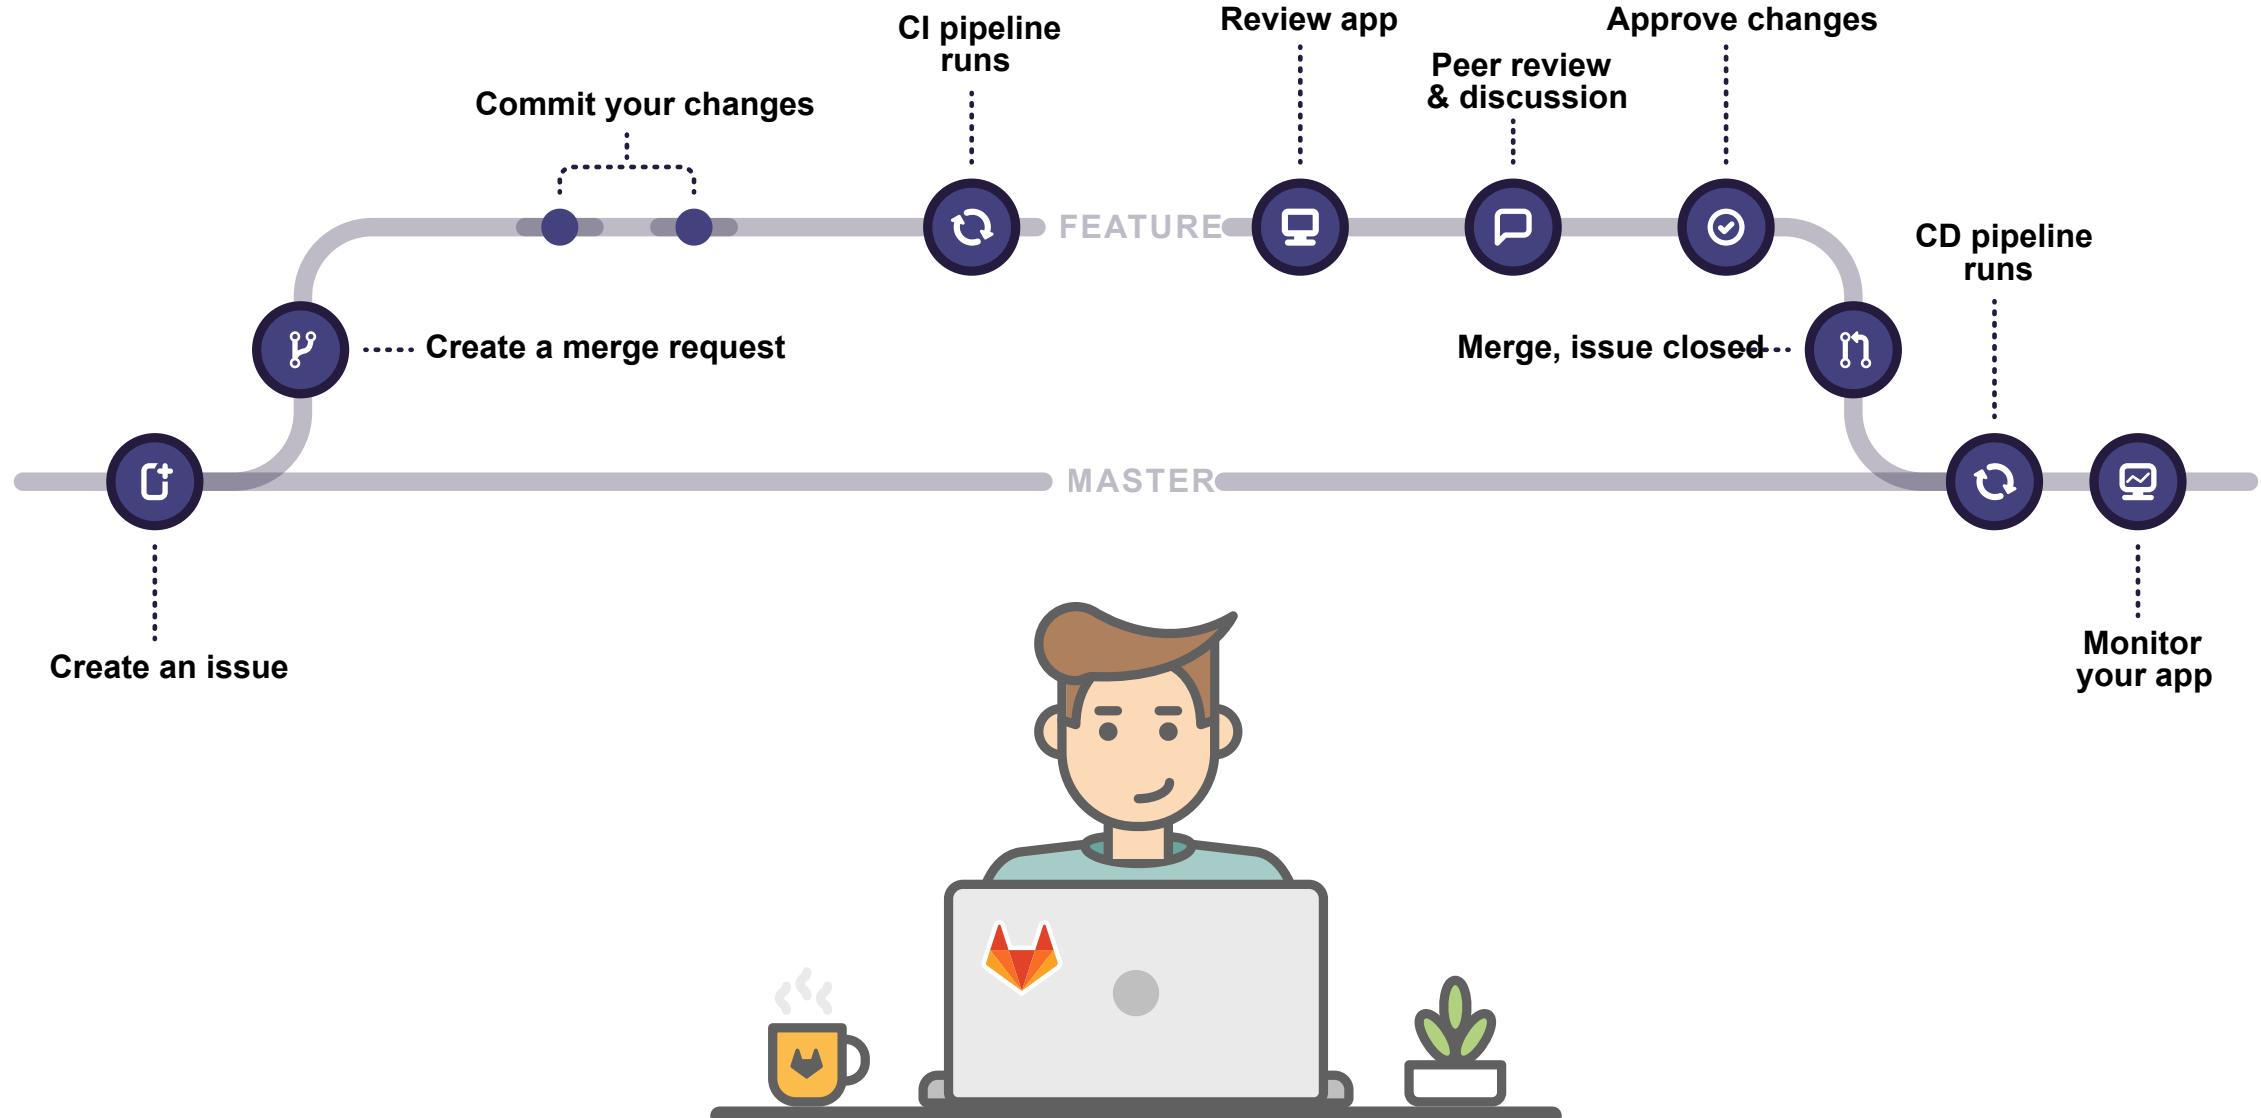
## Get your local up to date

- Use `pull` to update your local `master`

## Bonus

- Delete your branch & prune the remote

# This is the GitLab Flow



## Remote vs Local

Two copies of the repository are now at play

Our **local** keeps track of the remote's branches

The **remote** has no idea we exist...

# Collaborating in a repository

# Demo - Set up to work as a team

*Just watch while I...*

- Make a **public sub group** for our project(s)
- Add a new "About Us" Project w/ some branching rules
- Add some tasks to our board, maybe milestones
  - #1 "Add a Readme"
  - #2 "Add Ryan's Profile"
  - #3 "Add an index"
  - #4 "Add your name to the index"
  - #5 "Build HTML & Deploy"
- And finally complete task #1

*When I'm done, you should:*

*Request Access to the Sub Group!*

# Groups

- Assemble related projects in a "folder" like structure
- Give access to several projects at once
- Can be nested to match org structure
- Own Epics (along with boards & milestones)

We'll use a group to define a "team" with their own project(s) that they can access

# Namespaces in GitLab

Projects will exist in group or user namespaces

```
http://gitlab.example.com/username  
http://gitlab.example.com/groupname  
http://gitlab.example.com/groupname/subgroup_name
```

# Issue Management

- Epics
- Milestones (ie: sprints)
- Issues
- Labels
- Boards

Some features are only available at both a project and group level.

# Cloning

- Copy repository to another location
- Sets up `master` branch and `origin` remote
- You can choose a folder if you prefer

```
git clone git@gitlab.com:user/repo_name  
cd repo_name
```

# Lab: Add & Resolve an Issue

- Make a new issue in the **About Us** Project
  - Add my <name> Profile
  - Assign it to yourself
- Go local and complete the issue
  - clone the repository
  - Create a branch from master
    - Name it after the issue #
  - Do your work and push it up
- Create a Merge Request
  - Request to merge your branch into master
- Go *review* your teammates' Merge Requests

*Don't merge anything*

## Lab: Update your branch

Pretend I reviewed your code and requested a small change, ex: add your favorite animal or clean up some formatting.

- Make the update locally
- Push it up to the same branch
- Refresh your merge request

# Keeping branches up to date

- Get the latest master
  - Merge the remote `master` into your topic branch
- Get updates from the topic branch itself
  - Merge the remote topic branch into your local topic branch

Demo: I'll add an `index.md` file if I haven't already

# Lab: Update & Collaborate

The task: Add your name to the `index.md` file

- Make sure your local `master` is up to date

```
git checkout master  
git pull
```

- Then do the work
  - Branch off master
- Finally, share your work
  - Push it up and open a merge/pull request

*Don't merge it*

## Conflicts with Remotes

When you get a conflict in a merge/pull request, you'll need to make the fix locally

On your local "topic" branch (that is conflicted):

- `merge` in the latest copy of `master`
  - this should cause the conflict
- resolve the conflict
- push up the update

# Lab: Dealing with conflicts

The task: Fix the conflict and update your merge/pull request

- Make sure `master` is up to date
- Update your branch to cause the conflict
  - Merge `master` into your branch
  - or simply `pull origin master` into your branch
- Resolve the conflict
  - Fix, stage, commit
- Share
  - Push it up!
  - View your updated merge/pull request

## Forks

- Break one project down into many sub-projects
- Copy another team's project so you can use it or contribute to it
- Typical of open source contributors
- GitLab/GitHub thing, not a git thing

## Typical fork flow

- Personal fork will act as a go-between to the "main" project

```
Main Repo    -> Fork to my account -> Clone to local  
(upstream)          (my origin)
```

- push to your personal origin
- Merge Request into the main repo
- pull from the main repo, the upstream

**module**

# GitLab Pages

# GitLab Pages

Host a static site on GitLab:

```
<username or groupname>.gitlab.io/<project name>
```

- Lots of templates out there
- Anything in a `./public` folder will be exposed
  - If you build this folder you'll need to export as an `artifact`
- Requires GitLab CI & a `job` labeled `pages`

# Getting started with Pages

Create a `.gitlab-ci.yml` file

```
image: alpine:latest

pages:
  stage: deploy
  script:
    - echo 'Nothing to do...'
artifacts:
  paths:
    - public
only:
  - master
```

# CI/CD in GitLab

# Modern Software Lifecycle

We want to plan, implement and release to market as quickly as possible, then collect feedback, rinse and repeat.

And to do it with:

- Speed
- Agility
- Stability
- Flexibility

*How do we get there?*

## At the DevOps level

- Tightly integrate dev & operations teams
- Automate *All The Things*
- Keep configuration close to code
- Treat infrastructure like code

## And at the feature/code level

- Make small units of changes
- Which are easier to review, test and release
- Merge quickly to avoid conflicts
- Reduce risk and increase speed to market

# Continuously

The faster we can release, the more feedback we get and the better we can respond.

- Integrate
- Build, Test, Analyze
- Deploy

There are a lot of things we may want to do:

- Write code
- Run tests
- Check code quality
- Check for security issues
- Containerize
- Build
- Deploy
- Monitor

And we want to do it *frequently* and *quickly*

**GitLab wants to be the central platform for your entire software development pipeline from code to deploy & monitoring**

Unify your development and operations teams on one platform.

Reduce and remove "brittle" connections between many disparate platforms/services.

"spend more time writing code, less time maintaining the tool chain"

One end-to-end tool

## We'll need a few additional tools to pull this off in GitLab

- YAML (for configs)
  - And knowledge of their pipelines
- Basic knowledge of containers (ie: docker)
- Going cloud native (containerize) will make this easier

# Containerization

- *Virtualized* instances of your application bundled with all dependencies and services
- Cloud-native (but doesn't have to be)
- Docker & Kubernetes

# YAML

YAML Ain't Markup Language

- Flexible and data-oriented
- Indentation matters!
- YAML primer [here](#) and [here](#)

# YAML Example

```
# comment line
# use 2 spaces to indent
a_nested_map:
  key: value
  another_key: value

# inline sequence
[milk, eggs, juice]

a_sequence_or_array:
  - Item 1
  - Item 2
  -
    - Another
    - sequence indice

data: |
  Block of content
  One new lines
```

# GitLab Pipelines

Just requires a `.gitlab-ci.yml` file to be present.

Broken down into *stages* & *jobs*

- Stages
  - Logical grouping of work
  - ie: build, test, deploy
- Jobs
  - Each stage can contain one or more
  - Can pass files between stages

## Demo: CI/CD with Pipelines

*I'll set up a basic pipeline while we learn some of the configurations*

I can force a failure: `test -f README.md && return 1 || return 0`

# Stages

- `stages` map defines order of stages
  - defaults to `build`, `test` and `deploy`
- The default `stage` for any *job* is `test`
- You can skip with `.`
- `.pre` and `.post` are always first/last

```
stages:  
- build  
- test  
- deploy  
  
magical_flower:  
  stage: build  
  script: "echo 1"
```

# Jobs

- Run in parallel
- When all jobs are complete, next stage begins
- A job will fail if any non-0 value is returned
  - But jobs can be allowed to fail
- Must contain `script` clause
- No limit to the # of jobs
- Read up on [configuration parameters](#)

## Example jobs

```
job1:  
  stage: test  
  script: "./build.sh"  
  allow_failure: true  
  
test:  
  script:  
    - echo "OK, Going to test"  
    - npm install  
    - npm test
```

# Pipeline Images

Can define a container image for the entire pipeline, or per job.

```
image: alpine

test:
  image: node:latest
  script:
    - echo "OK, Going to test"
    - npm install
    - npm test
```

Need custom dockerfiles each build? [Use the docker-in-docker image](#)

# Control *when* a job or stage runs

`when` controls when the job runs, such as `manual`, `delayed`, etc...

```
cleanup_job:
  stage: cleanup
  script:
    - ./cleanup.sh
  when: on_failure

deploy:
  stage: deploy
  environment: production
  when: manual
  script:
    - apt-get update -qy
    - apt-get install -y ruby-dev
    - gem install dpl
    - dpl --provider=heroku --api-key=$HEROKU_API_KEY --app=gitlab-demo-ci
```

# Caching between stages

```
cache:  
  key: ${CI_COMMIT_REF_SLUG}  
  paths:  
    - node_modules/  
  
before_script:  
  - yarn install
```

# Artifacts

- You can pass data between Stages (not jobs)
- Can be set to expire
- Can be downloaded or pushed into reporting tools
- [Read more in the docs](#)

```
build:  
  stage: build  
  script: make build  
  artifacts:  
    - build  
    - .config
```

[Read more in the docs](#)

# Only / Except

- Can run jobs or stages *only* on specific branches
  - `only` defines when the job will run
  - `except` defines when it will *not* run
- They are inclusive (you can use both together)
- Default `only` is `['branches', 'tags']`

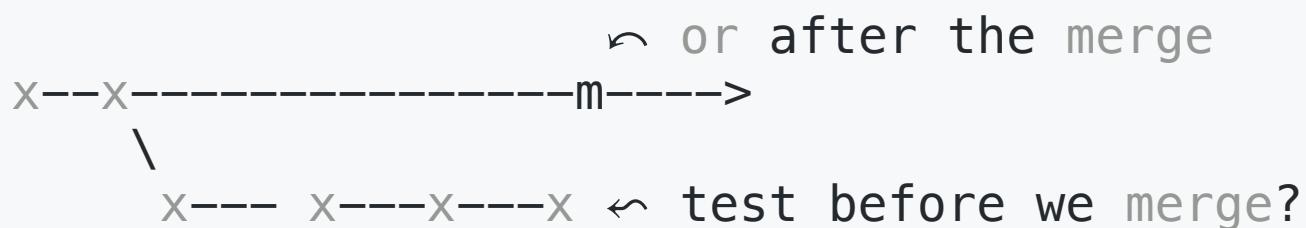
```
review:  
only:  
  - master
```

Warning: may be deprecated in favor of `rules`

# Pipelines for Merge Requests

Protected against "race condition" conflicts

- Set `only: merge_requests` and it will detach the merge request and test it (not testing the result)
- In combo with `Attempt to test the merge result` it will test the merge between the MR and master



## Merge Trains

Ability to queue a chain of merge requests, each built against the previous merge

- Requires "pipelines for merge results" enabled
- Queues merges but *immediately* begins their pipelines on the result
  - This can be wasteful if you skip the queue
- Each merge will be done against the previous branch in the chain
- Max of 20 will be run in parallel

[Read more](#)

## Environment Vars in the Pipeline

These can be defined per project in the UI

As well as within the pipelines for all stages, or individual stages and jobs.

```
variables:  
TEST: "HELLO WORLD"
```

# Deploying

- You can deploy anywhere!
  - S3
  - EC2 on AWS
  - Heroku
  - Google Cloud
  - Azure
  - bare metal...
- Tight integration with Kubernetes
- Use [DPL](#) as your deployment cli tool

```
staging:  
  stage: deploy  
  script:  
    - gem install dpl  
    - dpl --provider=heroku --app=my-app-staging --api-key=$HEROKU_STAGING_API_KEY
```

# Deploy to AWS

- Use the AWS image to get cli commands
- Add your AWS key/secret to your env vars
- Specify the image for your deployment

```
deploy:  
  stage: deploy  
  image: registry.gitlab.com/gitlab-org/cloud-deploy:latest # see the note below  
  script:  
    - aws s3 ...  
    - aws create-deployment ...
```

[Read the docs](#)

# Deploy to Google App Engine

- Create a service account role in google App engine
- Create a JSON key
- Add PROJECT\_ID & SERVICE\_ACCOUNT env vars

```
image: google/cloud-sdk:alpine

deploy_production:
  stage: deploy
  environment: Production
  only:
    - master
  script:
    - echo $SERVICE_ACCOUNT > /tmp/$CI_PIPELINE_ID.json
    - gcloud auth activate-service-account --key-file /tmp/$CI_PIPELINE_ID.json
    - gcloud --quiet --project $PROJECT_ID app deploy app.yaml dispatch.yaml
```

(Read the docs) [<https://medium.com/google-cloud/automatically-deploy-to-google-app-engine-with-gitlab-ci-d1c7237cbe11>]

## Demo: Team Pipeline

Let's get a pipeline going for our team project

- I want to build html from our markdown
- I want to deploy it to "GitLab Pages"

*These have working builds & pipelines I could steal from:*

<https://gitlab.com/rm-training/fork-me-for-the-team>

<https://gitlab.com/rm-training/demo-about-me>

# Lab: Basic Pipeline Playground

## Create a basic pipeline

- Start a new Project
- Create the `.gitlab-ci.yml` file
- Set up 3 `stages` with at least 5 `jobs`
- Pass an `index.html` through each stage
- Have it "deploy to pages"
  - By putting html files in `./public`

## Bonus:

- Make the deploy step `manual`
- Add a merge request that fails the pipeline

# GitLab Runners

- Executes pipeline *Jobs*
- Open Source
- Scalable

# Runner Types

- Shared (on [GitLab.com](#))
  - but may be slow, unavailable and uses "minutes"
  - may want a more secure environment
- Local
- Self-hosted

# Runners

- You can install them locally or on another instance
- Runner tags can define runners that only handle specific jobs

**module**

# Auto DevOps

| Zero-config CI/CD

## Demo: Auto DevOps

I'll set up a new project from the Node/Express template and turn on Auto DevOps

- View the pipeline(s) created
- Make a merge request to edit the title of the site

I have a project that is ready to look at as an "end state" of the demo:

[Demo Auto DevOps with Node](#)

# Auto DevOps

- Pretty awesome... when your app fits into the requirements
  - Deploys automatically into Kubernetes clusters (GKE, Amazon)
- Can be customized & extended
- Uses Herokuish and buildpacks to detect and auto-build, test, etc...
- "Best Practices" built-in

- Auto Build
- Auto Test
- Auto Code Quality
- Auto SAST (Static Application Security Testing)
- Auto Dependency Scan
- Auto License Compliance Scan
- Auto Container Scanning
- Auto "Review Apps"
- Auto DAST (Dynamic...)
- Auto Browser Performance Testing
- Auto Deploy
- Auto Cleanup

[Read up](#)

# Extending Auto DevOps

Easy to customize and extend pipelines

Copy [the Auto DevOps template](#) as needed.

```
include:
  - template: Auto-DevOps.gitlab-ci.yml

build_merge:
  stage: build
  extends: build
  only:
    - merge_requests
```

You can also disable jobs via ENV vars, ex `CODE_QUALITY_DISABLED`

## Define your own Dockerfile

Auto DevOps can be limiting, if you have your own Dockerfile, [GitLab will use that](#).

But you may need the *Container Registry* enabled.

# Example Dockerfile

```
FROM node:8.11-alpine

WORKDIR /usr/src/app

ARG NODE_ENV
ENV NODE_ENV $NODE_ENV

COPY package.json /usr/src/app/
RUN npm install

COPY . /usr/src/app

ENV PORT 5000
EXPOSE $PORT
CMD [ "npm", "start" ]
```

# Deploying to Kubernetes Clusters

# Kubernetes

- Can use Docker, or other container systems
- Automates provisioning, load balancing, etc
- Manages a collection of *nodes*
- Cluster
  - Node(s) (Worker Machine)
    - Pod(s)